

Technical Report: AI-Powered Resume Ranking System

Developed by: Rakshitha S (AIML Intern)

Abstract

The Resume Ranker AI is a specialized application designed to automate and streamline the candidate screening process for recruiters and hiring managers. Leveraging advanced Natural Language Processing (NLP) techniques, the system calculates a semantic similarity score between job descriptions (JD) and applicant resumes. The primary goal is to provide an objective, data-driven ranking of candidates, thereby reducing the time spent on manual resume review. The entire system is packaged within a user-friendly web interface built with Gradio, allowing users to upload multiple documents, receive instantaneous scores and visualizations, and export the results for further analysis. This tool significantly enhances efficiency in high-volume recruitment scenarios.

Introduction

In modern recruitment, the sheer volume of applications often creates a bottleneck at the initial screening stage, leading to potential biases and inconsistencies. The Resume Ranker AI addresses this challenge by implementing an intelligent scoring mechanism. This project aims to replace subjective, time-consuming manual screening with an efficient, objective, and scalable AI-driven solution. The application is built around the core concept of semantic similarity, which moves beyond simple keyword matching to understand the actual meaning and relevance of a candidate's experience and skills in the context of the job description.

Tools Used

The project relies on a minimal yet powerful stack of Python libraries, all utilized within a single Jupyter Notebook environment (`resume_ai.ipynb`).

- **Gradio:** Used to create the user-friendly web interface, enabling drag-and-drop file uploads and interactive display of results without the need for complex front-end development.
- **Sentence-Transformers:** A specialized framework used for generating high-quality vector embeddings of the text content from the resumes and the JD. This allows the system to measure *semantic* (meaning-based) similarity.
- **Scikit-learn (e.g., `metrics.pairwise.cosine_similarity`):** Used to perform the core similarity calculation, comparing the vector embeddings of the resumes against the JD.
- **pdfplumber & python-docx:** Essential utility libraries for robustly extracting raw text content from the two supported file formats: PDF (`.pdf`) and Microsoft Word (`.docx`).
- **Pandas & Matplotlib:** Used for data structuring (dataframes) and generating visualizations (bar charts) of the final ranking scores.

Steps Involved in Building the Project

The development of the Resume Ranker AI followed a structured data science pipeline, from data ingestion to final presentation.

1. **Dependency Setup and Initialization:** The required libraries (Gradio, Sentence-Transformers, pdfplumber, etc.) were installed and imported. Utility functions for file parsing and data handling were defined.
2. **Document Parsing and Text Extraction:** A crucial initial step involved defining functions (`extract_text`) to reliably read and extract raw, clean text from both `.pdf` and `.docx` files using `pdfplumber` and `python-docx`. This ensured that all document types could be processed uniformly.
3. **JD and Resume Separation:** A mechanism was implemented to scan the uploaded files and heuristically separate the single Job Description (JD) file from the multiple candidate Resumes, enabling targeted processing.
4. **Semantic Vectorization:** The core text content of the JD and all resumes were converted into dense numerical vectors (embeddings) using a pre-trained `Sentence-Transformer` model. These embeddings mathematically represent the semantic meaning of the text.
5. **Similarity Scoring and Ranking:** The embeddings of each resume were compared against the JD embedding using the `Cosine Similarity` metric. The resulting scores (ranging from 0 to 1) were then used to create a ranked Pandas DataFrame.
6. **Visualization and Export:** `Matplotlib` was used to generate a bar chart visualizing the similarity scores for easy comparison. The final ranked data was prepared for download as a CSV file.
7. **Web Interface Integration:** The entire logic was encapsulated within a single master function (`score_resumes`) and connected to a `gr.Interface` object in Gradio. This created the simple front-end for users to interact with the system by uploading files and viewing the results in real time.

Conclusion

The Resume Ranker AI successfully demonstrates the application of modern NLP and machine learning techniques to a critical business process. By providing a quantifiable match score and a ranked output, the system effectively automates the initial screening phase, allowing recruiters to focus human effort on the top-tier candidates. The use of Gradio ensures broad accessibility and ease of use, making the tool practical for immediate deployment. Future work could involve incorporating skill extraction, non-textual data analysis (e.g., date of experience), and advanced feedback mechanisms to further refine candidate fit.